

Jhabriel Allen & Jamier Carey

Professor Sridhar Ganti

ITAI 2372

July 4 2026

Introduction

Social media recommendation algorithms are one of the most common examples of artificial intelligence that people interact with every day, even when they do not realize it. Platforms such as Instagram, Facebook, and YouTube use these systems to decide which posts, videos, reels, shorts, advertisements, and suggested accounts appear in front of each user. The goal is to show people content they are likely to enjoy or find useful. However, the impact is much bigger than entertainment.

These systems can shape what people learn, what they believe, what communities they join, and even how they view people who are different from them. For this case study, I chose social media recommendation algorithms because they are helpful and risky at the same time. They make it easier to discover information, learn new skills, find creators, and connect with people or ideas that may not have appeared naturally. At the same time, they can keep people in a bubble, push one-sided information, reinforce stereotypes, and use personal data in ways that users may not fully understand.

Overview of the AI System

Social media recommendation algorithms are AI systems that help platforms decide which content to show each user. These systems are commonly powered by machine learning, which means they study patterns in user behavior and use those patterns to make predictions. If a person watches cooking videos, saves recipes, or follows food creators, the platform may

recommend more food content. If a person clicks on political or fitness videos, the system may show more of that content because it predicts continued engagement. Meta explains that its Instagram Feed Recommendations AI system automatically determines which content appears in a user's feed and in what order. The system considers signals such as user activity, information about the post, and how likely the user is to take an action, such as liking, sharing, commenting, or spending time on the post (Meta, 2025). YouTube describes its recommendation system as a way to help users find videos they want to watch, especially on the homepage and in the "Up Next" panel (YouTube, 2021). At a high level, these systems are not randomly selecting content. They are predicting what each person is most likely to respond to based on past behavior.

The main problem these systems solve is information overload. There is too much content online for any person to search through manually. Recommendation algorithms narrow that content down and bring users videos, posts, and creators that may match their interests. That can be helpful because users may discover information they did not know existed. However, the same system that helps people discover content can also decide too much for them if users do not have enough control over what they are being shown.

Benefits and Positive Impact

One major benefit is efficiency. Instead of users spending hours searching for information, the platform can quickly bring relevant content to them. This can help with entertainment, education, business, and general learning. A student may find study tips, a business owner may discover marketing advice, and someone learning a new skill may find useful tutorials. In that way, recommendation algorithms can make learning and discovery faster. These systems also help creators and businesses reach the right audiences. Without recommendations, a small creator may have a harder time being discovered by people who would actually enjoy their content.

With recommendation systems, a person who posts helpful videos about fitness, entrepreneurship, cooking, or technology can be shown to users who are already interested in those topics.

This creates opportunities for businesses, educators, and content creators who may not have a large following yet. There are also societal benefits. Recommendation algorithms can help people find communities, support systems, and useful information they may not have found offline. Pew Research Center found that many people across different countries believe social media helps inform people and raise awareness, even though they also worry it can manipulate and divide people (Wike et al., 2022). That shows the complicated nature of these systems. They can provide access and awareness, but they can also influence people in ways that may not always be healthy.

Ethical Issues

The first major ethical issue is bias and fairness. Recommendation systems are trained on data, and that data is based on human behavior. If people engage with biased content, stereotypes, or one-sided information, the algorithm may continue recommending similar content. This can create a loop where the person keeps seeing content that confirms what they already believe. For example, if someone already has a negative stereotype about a certain group of people, the algorithm may recommend more content that supports that stereotype because it sees the person engaging with it. That does not mean the stereotype is true. It means the system is rewarding engagement, even when the content may be harmful or unfair.

The second ethical concern is privacy and data usage. Social media recommendation algorithms rely heavily on user data. They may consider likes, watch time, comments, shares, searches, follows, and other activity. Many users know that platforms track some behavior, but

they may not understand how much of their behavior is being used to shape what they see next. This creates a consent issue. A user may technically agree to platform terms, but that does not always mean they fully understand how their data is being used to influence their feed.

The third ethical issue is the ability of these systems to steer opinions. If a person shows interest in a political topic, the system may continue showing content from the same viewpoint. This can make it seem like most people agree with that viewpoint, even when the platform is only showing a narrow slice of reality.

Societal Implications

The societal implications are serious because these systems affect how people understand the world. One concern is the impact on public opinion. If users mainly see information that supports their existing beliefs, they may become less open to other perspectives. A systematic review on filter bubbles found that personalization can isolate people from diverse opinions and reinforce existing beliefs (Areeb et al., 2023). This connects directly to the concern that people may not be exposed to the full picture.

These systems can also affect inequality and discrimination. If algorithms keep recommending content that reinforces negative stereotypes, they can contribute to unfair treatment of certain groups. This matters because social media is not just for entertainment anymore. It is where people learn about politics, culture, health, jobs, and communities. If the content people see is biased, then their understanding of real people and real issues can become biased too.

There is also an impact on jobs and the workforce. Recommendation algorithms influence how creators, marketers, journalists, and small businesses get attention. This creates opportunity because someone can reach a large audience without traditional media access, but it also

creates instability when income and visibility depend on an algorithm people do not control or fully understand.

Public trust is another major issue. If users feel that platforms are manipulating them or hiding why certain content appears, trust decreases. Globally, these systems can also create cultural problems because a recommendation system designed around engagement may not fully understand local context, cultural sensitivity, or the social impact of misinformation.

Risks and Challenges

One major risk is that users may only see what the algorithm wants them to see. This does not always happen in an obvious way. A person may simply scroll every day, not realizing that the system is slowly narrowing what they are exposed to. Over time, this can affect how they think, what they believe, and how they view other people. Users may feel like they are making independent choices, but the platform is heavily shaping the menu of choices in front of them.

A technical challenge is that recommendation systems are difficult to balance. Platforms want to show users content they will engage with, but engagement is not the same as truth, fairness, or quality. Sometimes content that is emotional, extreme, or controversial gets more attention. If the system rewards engagement too much, it may push content that keeps people watching but does not help them become better informed. Research on recommender systems has identified fairness, bias, threats, and privacy as major challenges because these systems can unintentionally reinforce inequalities and create security concerns (Roy et al., 2024).

Recommendations

A practical recommendation is to give users more direct control over recommendation settings. Users should be able to turn personalized recommendations on or off. They should also be able

to reset their recommendation history, choose whether they want more diverse viewpoints, and clearly tell the system when they do not want certain types of content. Some platforms already offer pieces of this, but it needs to be easier, clearer, and more consistent.

Platforms should also explain why content is being recommended. For example, a post could say, “You are seeing this because you watched similar videos,” or “You are seeing this because you follow accounts related to this topic.” This would help users understand the connection between their behavior and their feed. It would also help users make more informed choices about whether they want to keep engaging with that type of content.

Another recommendation is stronger bias testing. Social media companies should regularly audit recommendation systems to see whether they are reinforcing stereotypes, misinformation, discrimination, or one-sided content. These audits should not only be internal. Independent researchers and regulators should have safe ways to evaluate platform systems without exposing private user data.

Policy and regulation should also play a role. Governments should consider rules requiring major platforms to be more transparent about how recommendation systems work, what data is used, and what controls are available. These rules do not have to reveal every company secret, but users should not be completely in the dark about systems that influence what they see every day.

Conclusion

Social media recommendation algorithms are powerful AI systems that influence what people see, learn, and believe every day. They solve a real problem by helping users sort through massive amounts of online content and discover posts, videos, creators, and communities that

match their interests. They also create benefits for businesses, educators, creators, and everyday users by making information easier to find and personalize.

At the same time, these systems come with serious ethical concerns. They can reinforce bias, use large amounts of personal data, operate without enough transparency, and push people into content bubbles where they mainly see what confirms their existing beliefs. The issue is not that recommendations are automatically bad. The issue is that users often do not have enough knowledge or control over how recommendations shape their experience.

Citations

Areeb, Q. M., Nadeem, M., Sohail, S. S., Imam, R., Doctor, F., Himeur, Y., Hussain, A., & Amira, A. (2023). Filter bubbles in recommender systems: Fact or fallacy: A systematic review. arXiv.

<https://arxiv.org/abs/2307.01221>

Meta. (2023, June 29). How AI influences what you see on Facebook and Instagram. Meta.

<https://about.fb.com/news/2023/06/how-ai-ranks-content-on-facebook-and-instagram/>

Meta. (2025, September 4). Instagram Feed Recommendations AI system. Meta Transparency Center. <https://transparency.meta.com/features/explaining-ranking/ig-feed-recommendations/>

Roy, F., Ding, X., Choo, K.-K. R., & Zhou, P. (2024). A deep dive into fairness, bias, threats, and privacy in recommender systems: Insights and future research. arXiv.

<https://arxiv.org/abs/2409.12651>

Wike, R., Silver, L., Fetterolf, J., Huang, C., Austin, S., Clancy, L., & Gubbala, S. (2022, December 6). Social media seen as mostly good for democracy across many nations, but U.S. is a major outlier. Pew Research Center.

<https://www.pewresearch.org/global/2022/12/06/social-media-seen-as-mostly-good-for-democracy-across-many-nations-but-u-s-is-a-major-outlier/>

YouTube. (2021, September 15). On YouTube's recommendation system. YouTube Official Blog.

<https://blog.youtube/inside-youtube/on-youtubes-recommendation-system/>